

# untitled4

November 13, 2025

```
[3]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score

df = pd.read_csv('/content/housing[1].csv')

display(df.head())
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	\
0	-122.23	37.88	41.0	880.0	129.0	
1	-122.22	37.86	21.0	7099.0	1106.0	
2	-122.24	37.85	52.0	1467.0	190.0	
3	-122.25	37.85	52.0	1274.0	235.0	
4	-122.25	37.85	52.0	1627.0	280.0	

	population	households	median_income	median_house_value	ocean_proximity
0	322.0	126.0	8.3252	452600.0	NEAR BAY
1	2401.0	1138.0	8.3014	358500.0	NEAR BAY
2	496.0	177.0	7.2574	352100.0	NEAR BAY
3	558.0	219.0	5.6431	341300.0	NEAR BAY
4	565.0	259.0	3.8462	342200.0	NEAR BAY

```
[2]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
Data columns (total 10 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   longitude        20640 non-null   float64
 1   latitude         20640 non-null   float64
 2   housing_median_age 20640 non-null   float64
 3   total_rooms      20640 non-null   float64
 4   total_bedrooms   20433 non-null   float64
```

```
5    population            20640 non-null  float64
6    households            20640 non-null  float64
7    median_income         20640 non-null  float64
8    median_house_value   20640 non-null  float64
9    ocean_proximity      20640 non-null  object
dtypes: float64(9), object(1)
memory usage: 1.6+ MB
```

```
[8]: print(df.isnull().sum())
```

```
longitude          0
latitude           0
housing_median_age 0
total_rooms        0
total_bedrooms    207
population         0
households         0
median_income      0
median_house_value 0
ocean_proximity    0
dtype: int64
```

```
[29]: df_cleaned = df.dropna()
print(f"Original DataFrame shape: {df.shape}")
print(f"DataFrame shape after dropping nulls: {df_cleaned.shape}")

display(df_cleaned.info())
```

```
Original DataFrame shape: (20640, 14)
DataFrame shape after dropping nulls: (20640, 14)
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
Data columns (total 14 columns):
 #   Column                Non-Null Count  Dtype  
--- 
 0   longitude             20640 non-null  float64 
 1   latitude              20640 non-null  float64 
 2   housing_median_age    20640 non-null  float64 
 3   total_rooms            20640 non-null  float64 
 4   total_bedrooms         20640 non-null  float64 
 5   population             20640 non-null  float64 
 6   households             20640 non-null  float64 
 7   median_income          20640 non-null  float64 
 8   median_house_value     20640 non-null  float64 
 9   ocean_proximity_<1H OCEAN 20640 non-null  bool   
 10  ocean_proximity_INLAND 20640 non-null  bool   
 11  ocean_proximity_ISLAND 20640 non-null  bool   
 12  ocean_proximity_NEAR BAY 20640 non-null  bool
```

```
13 ocean_proximity_NEAR OCEAN 20640 non-null bool  
dtypes: bool(5), float64(9)  
memory usage: 1.5 MB
```

None

```
[9]: print(f"Number of duplicate rows: {df.duplicated().sum()}"")
```

Number of duplicate rows: 0


```
[4]: X = df[['median_income']]  
y = df['median_house_value']  
  
if X.isnull().sum().any() or y.isnull().sum().any():  
    print("Warning: Missing values detected. Dropping rows with missing data  
    ↪for selected feature/target.")  
    combined = pd.concat([X, y], axis=1).dropna()  
    X = combined[['median_income']]  
    y = combined['median_house_value']  
  
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,  
    ↪random_state=42)  
  
print(f"Training data size: {len(X_train)} samples")  
print(f"Testing data size: {len(X_test)} samples")
```

Training data size: 16512 samples  
Testing data size: 4128 samples


```
[5]: model = LinearRegression()  
model.fit(X_train, y_train)  
  
print(f"Model trained successfully. Coefficients: {model.coef_[0]:.2f},  
    ↪Intercept: {model.intercept_:.2f}")
```

Model trained successfully. Coefficients: 41933.85, Intercept: 44459.73

```
[10]: plt.figure(figsize=(10, 6))  
sns.histplot(df['median_house_value'], bins=50, kde=True)  
plt.title('Distribution of Median House Value')  
plt.xlabel('Median House Value')  
plt.ylabel('Frequency')  
plt.grid(True)  
plt.show()
```



```
[11]: plt.figure(figsize=(10, 6))
sns.histplot(df['median_income'], bins=50, kde=True)
plt.title('Distribution of Median Income')
plt.xlabel('Median Income')
plt.ylabel('Frequency')
plt.grid(True)
plt.show()
```



```
[12]: residuals = y_test - y_pred

plt.figure(figsize=(10, 6))
sns.scatterplot(x=y_pred, y=residuals, alpha=0.6)
plt.axhline(y=0, color='red', linestyle='--', label='Zero Residuals')
plt.title('Residuals Plot')
plt.xlabel('Predicted Values')
plt.ylabel('Residuals')
plt.legend()
plt.grid(True)
plt.show()
```



```
[13]: numerical_df = df.select_dtypes(include=np.number)
correlation_matrix = numerical_df.corr()
display(correlation_matrix.head())
```

	longitude	latitude	housing_median_age	total_rooms	\
longitude	1.000000	-0.924664		-0.108197	0.044568
latitude	-0.924664	1.000000		0.011173	-0.036100
housing_median_age	-0.108197	0.011173		1.000000	-0.361262
total_rooms	0.044568	-0.036100		-0.361262	1.000000
total_bedrooms	0.069608	-0.066983		-0.320451	0.930380


  

	total_bedrooms	population	households	median_income	\
longitude	0.069608	0.099773	0.055310	-0.015176	
latitude	-0.066983	-0.108785	-0.071035	-0.079809	
housing_median_age	-0.320451	-0.296244	-0.302916	-0.119034	
total_rooms	0.930380	0.857126	0.918484	0.198050	
total_bedrooms	1.000000	0.877747	0.979728	-0.007723	


  

	median_house_value
longitude	-0.045967
latitude	-0.144160
housing_median_age	0.105623
total_rooms	0.134153
total_bedrooms	0.049686

```
[14]: plt.figure(figsize=(12, 10))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt=".2f",
             linewidths=.5)
plt.title('Correlation Matrix of Numerical Features')
plt.show()
```



```
[15]: selected_features = ['median_income', 'median_house_value',
                         'housing_median_age', 'total_rooms']
sns.pairplot(df[selected_features])
plt.suptitle('Pair Plot of Key Numerical Features', y=1.02)
```



```
[6]: y_pred = model.predict(X_test)

mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

print(f"Mean Squared Error: {mse:.2f}")
print(f"R-squared: {r2:.2f}")
```

Mean Squared Error: 7091157771.77

R-squared: 0.46

```
[7]: plt.figure(figsize=(10, 6))
sns.scatterplot(x=X_test['median_income'], y=y_test, alpha=0.6, label='Actual Values')
plt.plot(X_test['median_income'], y_pred, color='red', linewidth=2, label='Regression Line')
plt.title('Simple Linear Regression: Median Income vs. Median House Value')
plt.xlabel('Median Income')
plt.ylabel('Median House Value')
plt.legend()
plt.grid(True)
plt.show()
```

